

Rev 15 Errata



Special thanks to these readers for helping find and remove errors: Ha Nguyen, Jeremie Morfin, Josh Barron, Jordan Kovar, Alexis Gomez, Christopher Maxwell, Alex Lupsasca, Alex Huang, Gabe Arthur, Sean Huckabee, Tjado Edzards, Mark Lancaster, Lena Hickson Long, Simo Sorce, Alex Petrescu, Jarrod Sinclair, Ryan Evans, ReDay Zarra, Lauren Steely, Sage Trudeau, Dan Rosenstein, David Vediner, Marcus Alenius, Ilia Perlov, samy kamkar, Albert Lee, Robert Blumofe, Bryce Moore, Daniel Burrow, Eric Norige, Anonymous, Kyle Stemen, Riley Beeson, Jund Pascal, Michael Williams, Patrick Werner, Jaroslaw Marciniak, Nico Waser, Christian Holteng, Leo Li

Location	Description
p. 20, Fig. 1.13	Last panel should read $y < -x - 1$, and line should move down, not up.
p. 24, Fig. 1.16	Step 1000 label should be "dog", not "cat".
p. 27, Fig. 1.19	Clarification — Figure 1.19 shows a two-layer network solving XOR, but activations have not yet been introduced. As written, a purely linear two-layer network collapses to a single linear boundary and cannot represent XOR. This is later addressed in Figure 1.30 (p. 35), but when first encountering Figure 1.19 it can appear that the network is achieving nonlinear separation without activations as the accompanying text.
p. 33	Caption by the W_1X_1 graph should read "constant", not "contant".
p. 35, Fig. 1.30	In the enlarged left green node, the bias is labeled $b_1^{(2)}$, but it should be $b_2^{(1)}$.
p. 64	Missing x_2 term in the f_2 equation.
p. 68, Fig. 2.3 caption	Figure 2.4 should be Figure 2.5.
p. 188–189, Fig. 5.3	In the first row, far right, the token "Gener" should read "Gene". In the second row, far right, "ative" should read "rative".
p. 191, footnote 3	"preyty" → "pretty".
p. 205, third paragraph	"correspond the concept" should read "correspond to the concept".
p. 207, Fig. 5.20	Vocabulary size of GPT-3 corrected from ~100k to 50,257.
p. 247, second paragraph	Figure 7.1 should be Figure 7.2.
p. 252–253, Fig. 7.9	The three bottom arrows leading across the fold should be shifted down by one to point to Layers 20, 21, and 26, instead of 19, 20, and 25.

Location	Description
p. 254, Fig. 7.10	Small update to numbers in the Layer 21 output matrix.
p. 260, first paragraph	Figure 7.12 should be Figure 7.13.
p. 261, first paragraph	Figure 7.12 should be Figure 7.13.
p. 271, Fig. 8.4	"atttention" → "attention".
p. 280, first paragraph	"an" should be "and". "kv" should be "KV".
p. 284, Fig. 8.15 caption	"we know that our scores are computed as $L_{UK} \cdot W_{UK}$ " should read "we know that our keys are computed as $L_{KV} \cdot W_{UK}$ ".
p. 294, Fig. 9.4 caption	Clarification — the full process for this video generation requires 50 steps; here we're showing the results after steps 1, 10, 30, and 50.
p. 295, fifth paragraph	"call a" should read "call to a".
p. 295, sixth paragraph	Figure 9.3 should be Figure 9.4.
p. 297, first paragraph	"allows them learn" should read "allows them to learn".
p. 299, Fig. 9.8	The difference vector should point towards $\hat{I}$, not I .
p. 300, first sentence	Repeated the word "to" after the "9.2 Diffusion" heading.
p. 303, Fig. 9.12	Missing panel letters.
p. 303, last paragraph	"of understanding of what" → "of understanding what".
p. 314, first paragraph	Figure 9.23 should be Figure 9.25.
p. 329, first paragraph	Figure 9.35 should be Figure 9.34.
p. 329, footnote	"succesfully" → "successfully".
p. 369, Exercise 4.9	Final row — all 8s should be 32s, and the final number of max theoretical regions should be 12,768,774,001.